

```
import pandas as pd

import numpy as np

df =
pd.read_csv(r'https://raw.githubusercontent.com/Sarvesh-S-Patil/Dataset/main/MPG.csv')

df.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	

	model_year	origin	name
0	70	usa	chevrolet chevelle malibu
1	70	usa	buick skylark 320
2	70	usa	plymouth satellite
3	70	usa	amc rebel sst
4	70	usa	ford torino

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   mpg                   398 non-null   float64
1   cylinders              398 non-null   int64
2   displacement           398 non-null   float64
3   horsepower              392 non-null   float64
4   weight                 398 non-null   int64
5   acceleration           398 non-null   float64
6   model_year             398 non-null   int64
7   origin                 398 non-null   object
8   name                   398 non-null   object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
df.describe()
```

	mpg	cylinders	displacement	horsepower	weight
count	398.000000	398.000000	398.000000	392.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623
std	7.815984	1.701004	104.269838	38.491160	846.841774

min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model_year
count	398.000000	398.000000
mean	15.568090	76.010050
std	2.757689	3.697627
min	8.000000	70.000000
25%	13.825000	73.000000
50%	15.500000	76.000000
75%	17.175000	79.000000
max	24.800000	82.000000

df.nunique()

mpg	129
cylinders	5
displacement	82
horsepower	93
weight	351
acceleration	95
model_year	13
origin	3
name	305

dtype: int64

df.corr()

	mpg	cylinders	displacement	horsepower	weight
mpg	1.000000	-0.775396	-0.804203	-0.778427	-0.831741
cylinders	-0.775396	1.000000	0.950721	0.842983	0.896017
displacement	-0.804203	0.950721	1.000000	0.897257	0.932824
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538
weight	-0.831741	0.896017	0.932824	0.864538	1.000000
acceleration	0.420289	-0.505419	-0.543684	-0.689196	-0.417457

```
model_year      0.579267  -0.348746      -0.370164   -0.416361  -0.306564
```

```
          acceleration  model_year
mpg          0.420289      0.579267
cylinders    -0.505419     -0.348746
displacement -0.543684     -0.370164
horsepower   -0.689196     -0.416361
weight       -0.417457     -0.306564
acceleration  1.000000      0.288137
model_year    0.288137      1.000000
```

```
df=df.dropna()
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 392 entries, 0 to 397
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	mpg	392 non-null	float64
1	cylinders	392 non-null	int64
2	displacement	392 non-null	float64
3	horsepower	392 non-null	float64
4	weight	392 non-null	int64
5	acceleration	392 non-null	float64
6	model_year	392 non-null	int64
7	origin	392 non-null	object
8	name	392 non-null	object

```
dtypes: float64(4), int64(3), object(2)
```


```
memory usage: 30.6+ KB
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
sns.pairplot(df, x_vars =  
['displacement', 'horsepower', 'weight', 'acceleration', 'mpg'], y_vars=  
['mpg'])
```

```
<seaborn.axisgrid.PairGrid at 0x7f61c65f73d0>
```



```

df.columns
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')

y = df['mpg']
y.shape
(392,)

X= df[['displacement', 'horsepower', 'weight', 'acceleration']]
X.shape
(392, 4)

from sklearn.preprocessing import StandardScaler
ss= StandardScaler()
X =ss.fit_transform(X)
pd.DataFrame(X).describe()

      0      1      2      3
count  3.920000e+02  3.920000e+02  3.920000e+02  3.920000e+02
mean   -2.537653e-16 -4.392745e-16  5.607759e-17  6.117555e-16
std     1.001278e+00  1.001278e+00  1.001278e+00  1.001278e+00
min    -1.209563e+00 -1.520975e+00 -1.608575e+00 -2.736983e+00
25%    -8.555316e-01 -7.665929e-01 -8.868535e-01 -6.410551e-01
50%    -4.153842e-01 -2.853488e-01 -2.052109e-01 -1.499869e-02
75%     7.782764e-01  5.600800e-01  7.510927e-01  5.384714e-01
max     2.493416e+00  3.265452e+00  2.549061e+00  3.360262e+00

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.3,random_state=2529)
X_train.shape,X_test.shape,y_train.shape,y_test.shape
((274, 4), (118, 4), (274,), (118,))

from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train,y_train)
LinearRegression()
y_pred = lr.predict(X_test)

```

```
y_pred.shape
(118,)
lr.intercept_
23.485738559737584
lr.coef_
array([-1.05767743, -1.68734727, -4.10787617, -0.11495177])
from sklearn.metrics import
mean_absolute_error,mean_absolute_percentage_error,r2_score
mean_absolute_error(y_test,y_pred)
3.3286968643244106
mean_absolute_percentage_error(y_test,y_pred)
0.14713035779536746
r2_score(y_test,y_pred)
0.7031250746717692
```